Zero-Retraining Task Transfer via Structural Task Abstraction

Mayuri Gade | SIES Graduate School of Technology, Navi Mumbai, India

The Core Problem

The Problem

Current AI systems often need to be retrained every time the task changes. This makes them:

- Slow - Hours of retraining for each new task
- Expensive - Requires significant computational resources
- Hard to Scale - Impractical for real-world deployment

Human Intelligence

Humans instantly combine known skills to solve new problems. They:

- Adapt Immediately - No retraining needed
- Reuse Knowledge - Apply existing skills creatively
- Scale Effortlessly - Handle unlimited task variations

Introducing STAN

Learn Task Structure

STAN learns the internal structure of tasks—how smaller subtasks connect, how they combine logically, and how they form step-by-step procedures.

Compose Instantly

Once structure is learned, STAN can generate the required behavior for a new task in a single forward pass, without updating any model parameters.

Solve Efficiently

STAN can handle new tasks instantly in about 12 milliseconds, without any retraining.

Three Key Innovations

01

Three Structure Types

STAN represents tasks using three components: a graph capturing dependencies between subtasks, a logical representation capturing how conditions combine, and a program describing execution steps.

02

Single Forward Pass Generation

Instead of retraining a model, the Structure-Conditioned Hypernetwork produces the required behavior in one forward pass by adjusting weights based on the task structure.

03

Composition for New Tasks

The Composition Engine converts new task descriptions into structural representations, then combines them with previously learned components to solve new tasks instantly.

System Architecture

Structure Discovery Module (SDM)

Analyzes task examples and breaks them down into three parts: how subtasks depend on each other using graphs, how conditions and rules combine logically, and how tasks can be executed step by step.

Structure-Conditioned Hypernetwork (SCH)

Takes the learned task structure and generates the parameters needed to solve that task in one forward pass, without retraining.

Composition Engine (CE)

Used at inference time to convert new task descriptions into structural representations, then combines them with previously learned components to solve new tasks instantly.

Visual Task Transfer Results

Taskonomy Dataset | Lower Error = Better

STAN reduces error by nearly half vs. the strongest baseline ($p < 0.001$, $d = 3.8$).

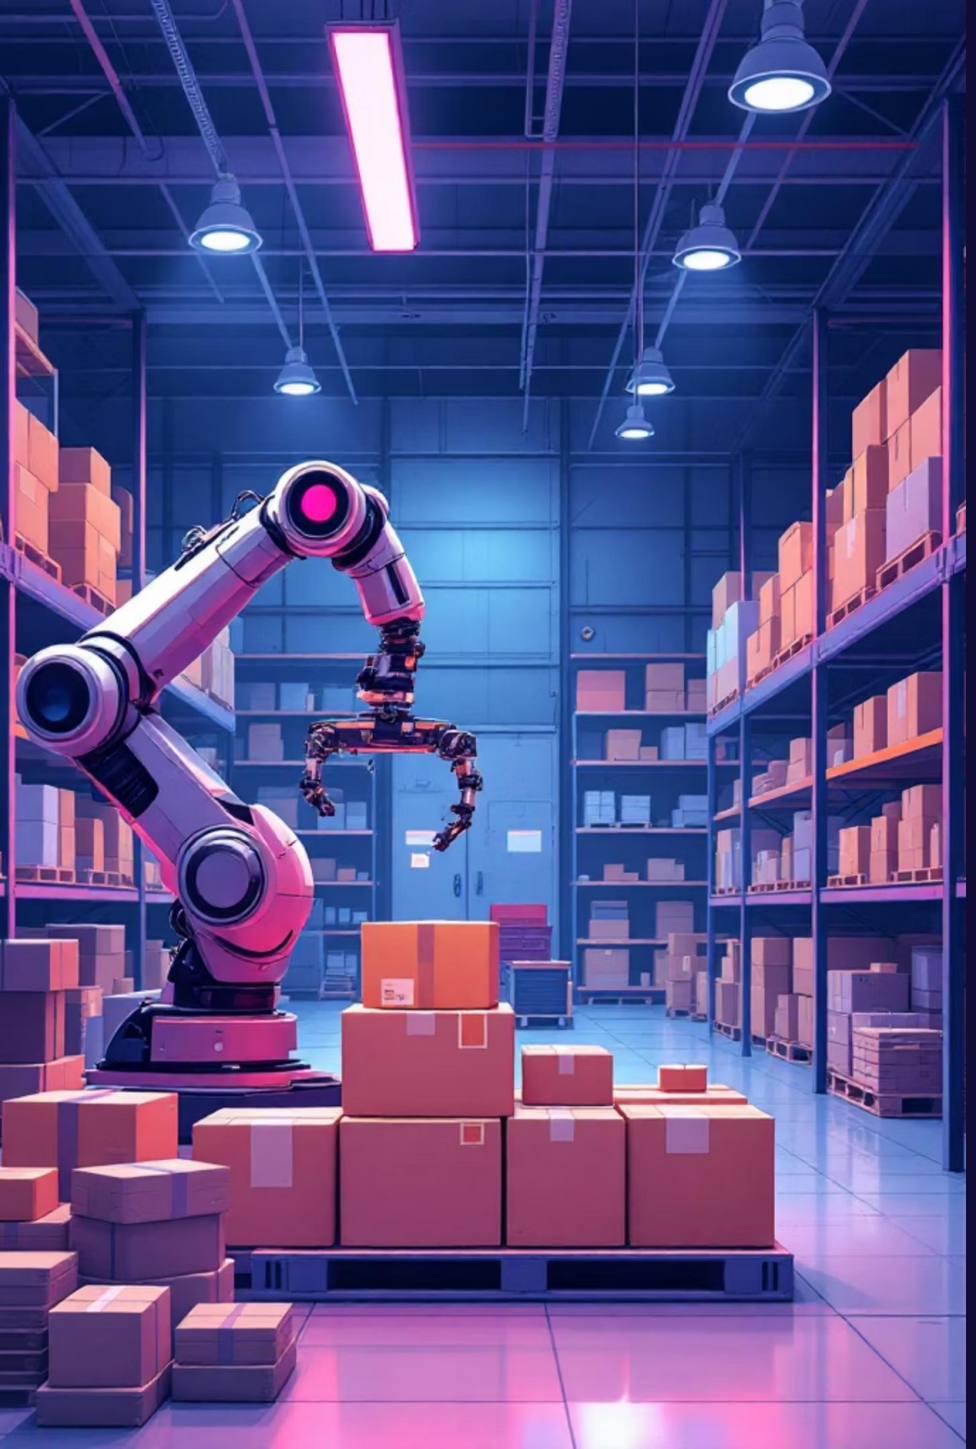

Robotics & RL Performance

Meta-World Robotics

Robotic Manipulation

89.4% Success Rate

STAN achieves close to 90% success rate, clearly outperforming all baselines. It handles logical conditions explicitly—understanding instructions like "pick non-fragile objects." This demonstrates real-world applicability.

Grid-World RL

Complex Task Composition

93.7% Success Rate

STAN successfully handles all 64 composed tasks, while other methods degrade significantly as complexity grows. It scales exceptionally well with structured and combinatorial tasks—proving its robustness.

Ablation Study: Component Importance

Success Rate (%) Across Task Depths

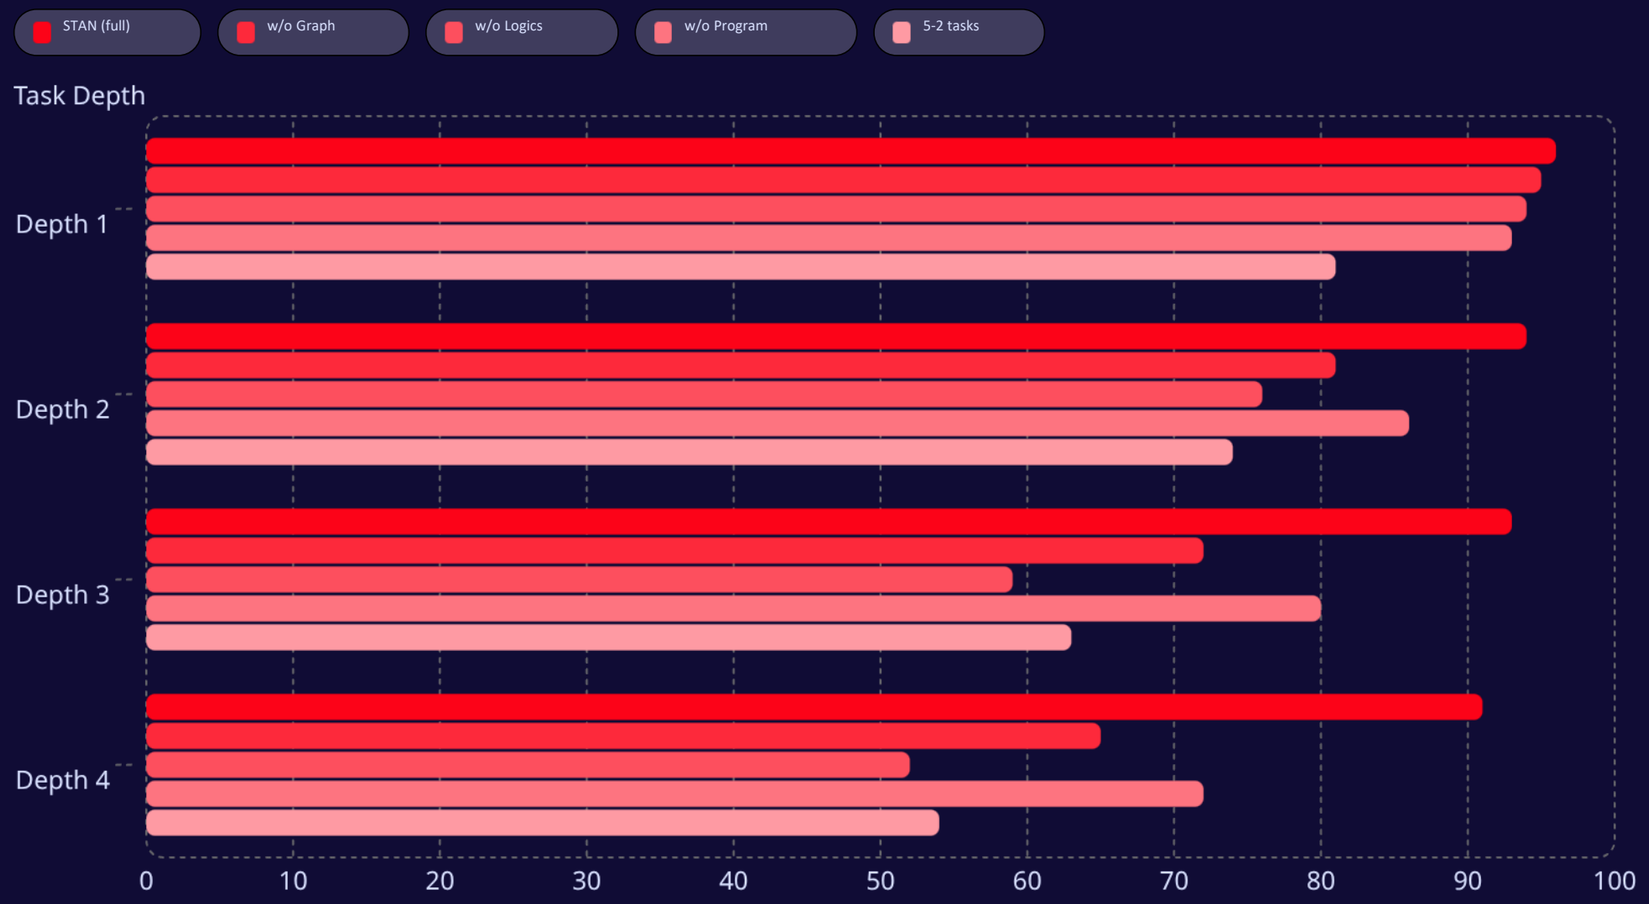

Each component plays a distinct, necessary role—especially for complex tasks.

Efficiency Gains

Speed Advantage

12ms vs 840ms

STAN solves new tasks in just 12 milliseconds. MAML takes 840 milliseconds. That's 70× faster—the difference between instant response and waiting almost a full second.

Computational Savings

97% Less Compute

Across 120 task transfers, STAN uses 97% less computational power than traditional retraining methods. This means lower energy costs, reduced carbon footprint, and faster execution on any device.

Instant Adaptation

Zero Retraining

Traditional methods require hours of retraining for each new task. STAN generates solutions in a single forward pass—instantly adapting without any additional training steps.

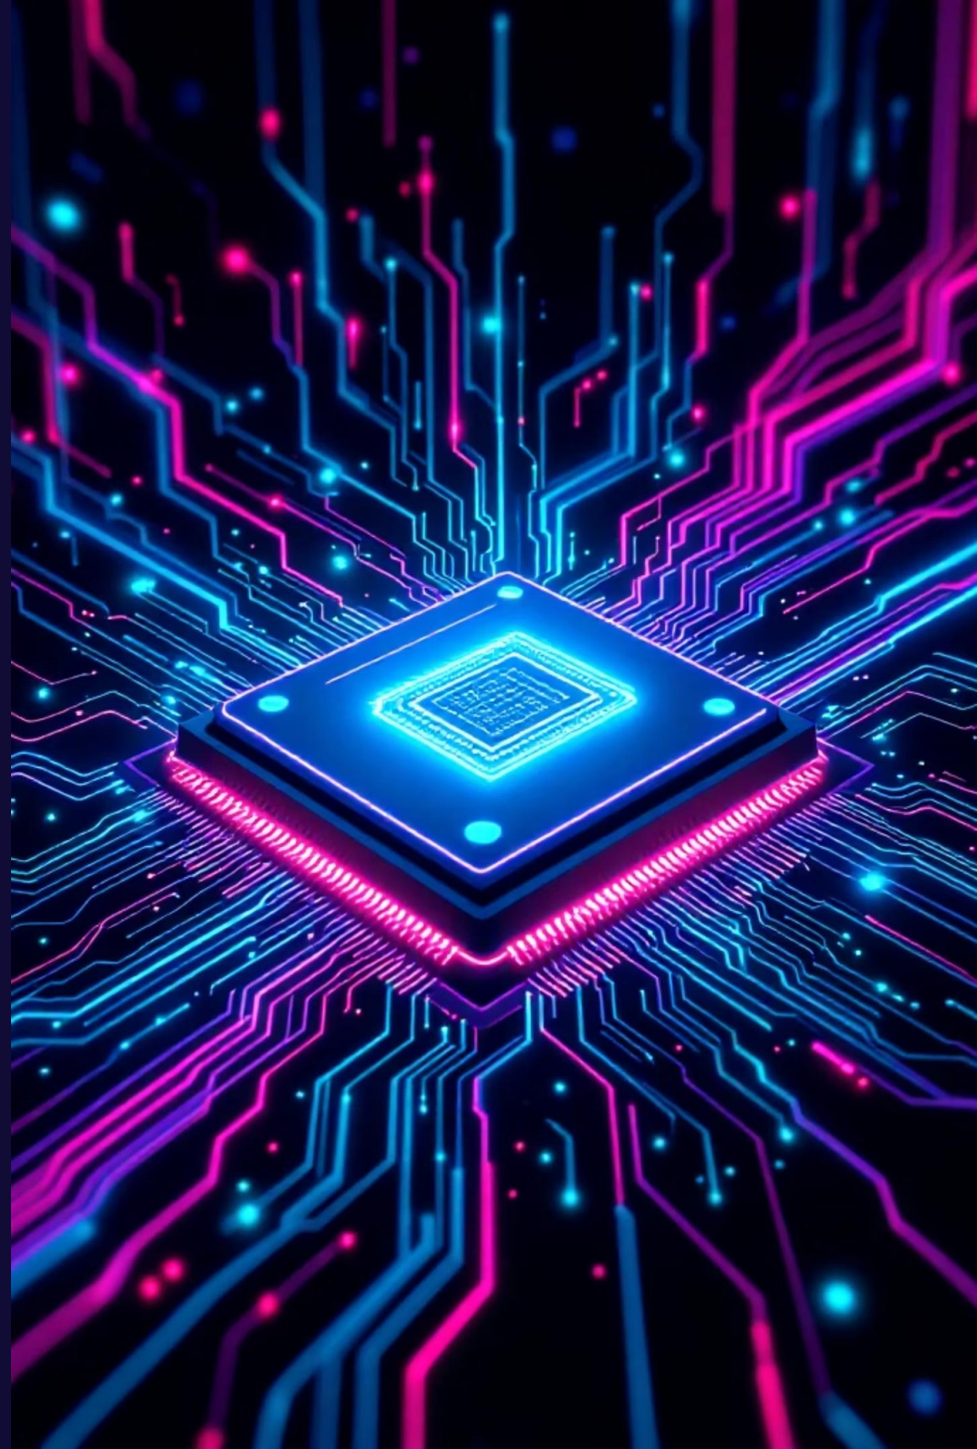

Future Directions & Conclusion

Expand Task Descriptions

The current system can handle a limited set of task descriptions. Expanding it to understand more open and flexible instructions—especially involving multiple agents—would make it far more practical.

Real-World Testing

So far, all experiments have been done in simulation. Testing the system on real robots is a crucial next step to evaluate how well it performs in real-world conditions.

Automatic Structure Learning

Right now, the system relies on some structured supervision. A key goal is to learn these structures automatically, without requiring labeled data.

Key Insight: Learning and reusing task structure enables faster, more efficient, and more flexible AI systems—moving closer to human-like generalization.